

# artificial\_neural\_network

June 21, 2021

## 1 Artificial Neural Network (ANN)

Work by: Raghavendra Tapas

### 1.0.1 Importing the libraries

```
[1]: import numpy as np
import pandas as pd
import tensorflow as tf
```

```
[2]: tf.__version__
```

```
[2]: '2.5.0'
```

### 1.1 Data Preprocessing

#### 1.1.1 Importing the dataset

- x - all the features except exited column
- y - value to be predicted

```
[22]: dataset = pd.read_csv('Churn_Modelling.csv')
x = dataset.iloc[:, 3:-1].values
y = dataset.iloc[:, -1].values
```

```
[23]: dataset
```

```
[23]:
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	\
0	1	15634602	Hargrave	619	France	Female	42	
1	2	15647311	Hill	608	Spain	Female	41	
2	3	15619304	Onio	502	France	Female	42	
3	4	15701354	Boni	699	France	Female	39	
4	5	15737888	Mitchell	850	Spain	Female	43	
...	...	...	...	...	...	...	...	
9995	9996	15606229	Obijiaku	771	France	Male	39	
9996	9997	15569892	Johnstone	516	France	Male	35	
9997	9998	15584532	Liu	709	France	Female	36	
9998	9999	15682355	Sabbatini	772	Germany	Male	42	
9999	10000	15628319	Walker	792	France	Female	28	

	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	\
0	2	0.00	1	1		1
1	1	83807.86	1	0		1
2	8	159660.80	3	1		0
3	1	0.00	2	0		0
4	2	125510.82	1	1		1
...	...	...	...	...	...	...
9995	5	0.00	2	1		0
9996	10	57369.61	1	1		1
9997	7	0.00	1	0		1
9998	3	75075.31	2	1		0
9999	4	130142.79	1	1		0

	EstimatedSalary	Exited
0	101348.88	1
1	112542.58	0
2	113931.57	1
3	93826.63	0
4	79084.10	0
...	...	...
9995	96270.64	0
9996	101699.77	0
9997	42085.58	1
9998	92888.52	1
9999	38190.78	0

[10000 rows x 14 columns]

```
[24]: print(x)
```

```
[[619 'France' 'Female' ... 1 1 101348.88]
 [608 'Spain' 'Female' ... 0 1 112542.58]
 [502 'France' 'Female' ... 1 0 113931.57]
 ...
 [709 'France' 'Female' ... 0 1 42085.58]
 [772 'Germany' 'Male' ... 1 0 92888.52]
 [792 'France' 'Female' ... 1 0 38190.78]]
```

```
[25]: print(y)
```

```
[1 0 1 ... 1 1 0]
```

### 1.1.2 Encoding Categorical Data

Label Encoding the “Gender” column

```
[26]: from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
x[:, 2] = le.fit_transform(x[:, 2])
```

```
[27]: type(x)
```

```
[27]: numpy.ndarray
```

One Hot Encoding the “Geography” column

```
[8]: from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import OneHotEncoder
ct = ColumnTransformer(transformers = [('encoder', OneHotEncoder(), [1])],
    ↳remainder='passthrough')
x = np.array(ct.fit_transform(x))
```

### 1.1.3 Splitting the dataset into the Training set and Test set

```
[9]: from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size = 0.2,
    ↳random_state = 0)
```

### 1.1.4 Feature Scaling

```
[10]: from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
x_train = sc.fit_transform(x_train)
x_test = sc.transform(x_test)
```

## 1.2 Building the ANN

How many neurons to choose is a trial and error process in deep learning.

- **1st layer:** rectifier activation function, 6 neurons
- **2nd layer:** rectifier activation function, 6 neurons
- **Output layer:** sigmoid activation function, 1 neuron, because the variable we are trying to predict is binary classification problem.

i.e. Whether or not the customer leaves the bank. (Yes or No)

### 1.2.1 Initializing the ANN

```
[28]: ann = tf.keras.models.Sequential()
```

### 1.2.2 Adding the input layer and the first hidden layer

```
[29]: ann.add(tf.keras.layers.Dense(units=6, activation='relu'))
```

### 1.2.3 Adding the second hidden layer

```
[30]: ann.add(tf.keras.layers.Dense(units=6, activation='relu'))
```

### 1.2.4 Adding the output layer

```
[31]: ann.add(tf.keras.layers.Dense(units=1, activation='sigmoid'))
```

## 1.3 Training the ANN

- **Adam optimizer:** Adaptive Moment Estimation or Adam optimizer combines two gradient descent methodologies. Adaptive Moment Estimation is an algorithm for optimization technique for gradient descent. The method is really efficient when working with large problem involving a lot of data or parameters.
- **Binary Cross-Entropy:** Cross-entropy is a measure of the difference between two probability distributions for a given random variable or set of events. Since we're trying to compute a loss, we need to penalize bad predictions. If the probability associated with the true class is 1.0, we need its loss to be zero. Conversely, if that probability is low, say, 0.01, we need its loss to be huge.

### 1.3.1 Compiling the ANN

```
[32]: ann.compile(optimizer = 'adam', loss = 'binary_crossentropy', metrics = ['accuracy'])
```

### 1.3.2 Training the ANN on the Training set

- Batch gradient descent is faster to compute. Usually batch size is taken to be 32.
- One epoch consists of one full training cycle on the training set. Once every sample in the set is seen, you start again - marking the beginning of the 2nd epoch.

```
[34]: ann.fit(x_train, y_train, batch_size = 32, epochs = 100)
```

```
Epoch 1/100
667/667 [=====] - 2s 3ms/step - loss: 0.3358 - accuracy: 0.8650
Epoch 2/100
667/667 [=====] - 2s 3ms/step - loss: 0.3339 - accuracy: 0.8644
Epoch 3/100
667/667 [=====] - 2s 3ms/step - loss: 0.3332 - accuracy: 0.8671
Epoch 4/100
667/667 [=====] - 2s 3ms/step - loss: 0.3335 - accuracy: 0.8636
Epoch 5/100
667/667 [=====] - 2s 3ms/step - loss: 0.3327 - accuracy: 0.8671
```

Epoch 6/100  
667/667 [=====] - 2s 3ms/step - loss: 0.3327 -  
accuracy: 0.8670  
Epoch 7/100  
667/667 [=====] - 2s 3ms/step - loss: 0.3319 -  
accuracy: 0.8666  
Epoch 8/100  
667/667 [=====] - 2s 3ms/step - loss: 0.3325 -  
accuracy: 0.8643  
Epoch 9/100  
667/667 [=====] - 2s 2ms/step - loss: 0.3323 -  
accuracy: 0.8639  
Epoch 10/100  
667/667 [=====] - 2s 2ms/step - loss: 0.3317 -  
accuracy: 0.8668  
Epoch 11/100  
667/667 [=====] - 2s 2ms/step - loss: 0.3321 -  
accuracy: 0.8648  
Epoch 12/100  
667/667 [=====] - 2s 3ms/step - loss: 0.3317 -  
accuracy: 0.8649  
Epoch 13/100  
667/667 [=====] - 2s 3ms/step - loss: 0.3319 -  
accuracy: 0.8629  
Epoch 14/100  
667/667 [=====] - 2s 3ms/step - loss: 0.3309 -  
accuracy: 0.8645  
Epoch 15/100  
667/667 [=====] - 2s 3ms/step - loss: 0.3310 -  
accuracy: 0.8677  
Epoch 16/100  
667/667 [=====] - 2s 3ms/step - loss: 0.3317 -  
accuracy: 0.8664  
Epoch 17/100  
667/667 [=====] - 2s 3ms/step - loss: 0.3309 -  
accuracy: 0.8649  
Epoch 18/100  
667/667 [=====] - 2s 3ms/step - loss: 0.3314 -  
accuracy: 0.8656  
Epoch 19/100  
667/667 [=====] - 2s 3ms/step - loss: 0.3314 -  
accuracy: 0.8665  
Epoch 20/100  
667/667 [=====] - 2s 3ms/step - loss: 0.3307 -  
accuracy: 0.8640  
Epoch 21/100  
667/667 [=====] - 2s 3ms/step - loss: 0.3308 -  
accuracy: 0.8662

Epoch 22/100  
667/667 [=====] - 2s 3ms/step - loss: 0.3304 -  
accuracy: 0.8654  
Epoch 23/100  
667/667 [=====] - 2s 3ms/step - loss: 0.3312 -  
accuracy: 0.8643  
Epoch 24/100  
667/667 [=====] - 2s 2ms/step - loss: 0.3306 -  
accuracy: 0.8656  
Epoch 25/100  
667/667 [=====] - 2s 3ms/step - loss: 0.3306 -  
accuracy: 0.8652  
Epoch 26/100  
667/667 [=====] - 2s 3ms/step - loss: 0.3301 -  
accuracy: 0.8654  
Epoch 27/100  
667/667 [=====] - 2s 3ms/step - loss: 0.3305 -  
accuracy: 0.8655  
Epoch 28/100  
667/667 [=====] - 2s 3ms/step - loss: 0.3304 -  
accuracy: 0.8652  
Epoch 29/100  
667/667 [=====] - 2s 2ms/step - loss: 0.3306 -  
accuracy: 0.8635  
Epoch 30/100  
667/667 [=====] - 2s 3ms/step - loss: 0.3298 -  
accuracy: 0.8662  
Epoch 31/100  
667/667 [=====] - 2s 3ms/step - loss: 0.3301 -  
accuracy: 0.8649  
Epoch 32/100  
667/667 [=====] - 2s 3ms/step - loss: 0.3292 -  
accuracy: 0.8652  
Epoch 33/100  
667/667 [=====] - 2s 3ms/step - loss: 0.3299 -  
accuracy: 0.8654  
Epoch 34/100  
667/667 [=====] - 2s 3ms/step - loss: 0.3293 -  
accuracy: 0.8652  
Epoch 35/100  
667/667 [=====] - 2s 3ms/step - loss: 0.3301 -  
accuracy: 0.8675  
Epoch 36/100  
667/667 [=====] - 2s 2ms/step - loss: 0.3295 -  
accuracy: 0.8654  
Epoch 37/100  
667/667 [=====] - 2s 2ms/step - loss: 0.3302 -  
accuracy: 0.8640

Epoch 38/100  
667/667 [=====] - 2s 2ms/step - loss: 0.3298 -  
accuracy: 0.8645  
Epoch 39/100  
667/667 [=====] - 2s 3ms/step - loss: 0.3296 -  
accuracy: 0.8627  
Epoch 40/100  
667/667 [=====] - 2s 3ms/step - loss: 0.3298 -  
accuracy: 0.8643  
Epoch 41/100  
667/667 [=====] - 2s 3ms/step - loss: 0.3296 -  
accuracy: 0.8640  
Epoch 42/100  
667/667 [=====] - 2s 3ms/step - loss: 0.3299 -  
accuracy: 0.8649  
Epoch 43/100  
667/667 [=====] - 2s 3ms/step - loss: 0.3298 -  
accuracy: 0.8675  
Epoch 44/100  
667/667 [=====] - 2s 3ms/step - loss: 0.3300 -  
accuracy: 0.8639  
Epoch 45/100  
667/667 [=====] - 2s 3ms/step - loss: 0.3286 -  
accuracy: 0.8650  
Epoch 46/100  
667/667 [=====] - 2s 2ms/step - loss: 0.3294 -  
accuracy: 0.8659  
Epoch 47/100  
667/667 [=====] - 2s 3ms/step - loss: 0.3285 -  
accuracy: 0.8652  
Epoch 48/100  
667/667 [=====] - 2s 3ms/step - loss: 0.3289 -  
accuracy: 0.8659  
Epoch 49/100  
667/667 [=====] - 2s 3ms/step - loss: 0.3294 -  
accuracy: 0.8650  
Epoch 50/100  
667/667 [=====] - 2s 3ms/step - loss: 0.3293 -  
accuracy: 0.8640  
Epoch 51/100  
667/667 [=====] - 2s 2ms/step - loss: 0.3289 -  
accuracy: 0.8643  
Epoch 52/100  
667/667 [=====] - 2s 3ms/step - loss: 0.3289 -  
accuracy: 0.8644: 0s - loss:  
Epoch 53/100  
667/667 [=====] - ETA: 0s - loss: 0.3291 - accuracy:  
0.8649 ETA: 0s - loss: 0.331 - 2s 3ms/step - loss: 0.3289 - accuracy: 0.8648

Epoch 54/100  
667/667 [=====] - 2s 3ms/step - loss: 0.3288 -  
accuracy: 0.8664  
Epoch 55/100  
667/667 [=====] - 2s 3ms/step - loss: 0.3291 -  
accuracy: 0.8656  
Epoch 56/100  
667/667 [=====] - 2s 3ms/step - loss: 0.3290 -  
accuracy: 0.8650  
Epoch 57/100  
667/667 [=====] - 2s 3ms/step - loss: 0.3283 -  
accuracy: 0.8650  
Epoch 58/100  
667/667 [=====] - 2s 3ms/step - loss: 0.3290 -  
accuracy: 0.8652  
Epoch 59/100  
667/667 [=====] - 2s 3ms/step - loss: 0.3285 -  
accuracy: 0.8658  
Epoch 60/100  
667/667 [=====] - 2s 2ms/step - loss: 0.3283 -  
accuracy: 0.8650  
Epoch 61/100  
667/667 [=====] - 2s 3ms/step - loss: 0.3287 -  
accuracy: 0.8654  
Epoch 62/100  
667/667 [=====] - 2s 2ms/step - loss: 0.3282 -  
accuracy: 0.8661  
Epoch 63/100  
667/667 [=====] - 2s 2ms/step - loss: 0.3283 -  
accuracy: 0.8659  
Epoch 64/100  
667/667 [=====] - 2s 2ms/step - loss: 0.3285 -  
accuracy: 0.8652  
Epoch 65/100  
667/667 [=====] - 2s 3ms/step - loss: 0.3283 -  
accuracy: 0.8650  
Epoch 66/100  
667/667 [=====] - 2s 3ms/step - loss: 0.3279 -  
accuracy: 0.8671  
Epoch 67/100  
667/667 [=====] - 2s 2ms/step - loss: 0.3280 -  
accuracy: 0.8630  
Epoch 68/100  
667/667 [=====] - 2s 3ms/step - loss: 0.3284 -  
accuracy: 0.8644  
Epoch 69/100  
667/667 [=====] - 2s 3ms/step - loss: 0.3279 -  
accuracy: 0.8675



Epoch 70/100  
667/667 [=====] - 2s 3ms/step - loss: 0.3283 -  
accuracy: 0.8648  
Epoch 71/100  
667/667 [=====] - 2s 3ms/step - loss: 0.3272 -  
accuracy: 0.8649  
Epoch 72/100  
667/667 [=====] - 2s 3ms/step - loss: 0.3281 -  
accuracy: 0.8656  
Epoch 73/100  
667/667 [=====] - 2s 3ms/step - loss: 0.3277 -  
accuracy: 0.8651  
Epoch 74/100  
667/667 [=====] - 2s 3ms/step - loss: 0.3279 -  
accuracy: 0.8668  
Epoch 75/100  
667/667 [=====] - 2s 3ms/step - loss: 0.3287 -  
accuracy: 0.8636  
Epoch 76/100  
667/667 [=====] - 2s 3ms/step - loss: 0.3280 -  
accuracy: 0.8658  
Epoch 77/100  
667/667 [=====] - 2s 3ms/step - loss: 0.3279 -  
accuracy: 0.8656  
Epoch 78/100  
667/667 [=====] - 2s 3ms/step - loss: 0.3279 -  
accuracy: 0.8656  
Epoch 79/100  
667/667 [=====] - 2s 3ms/step - loss: 0.3278 -  
accuracy: 0.8633  
Epoch 80/100  
667/667 [=====] - 2s 3ms/step - loss: 0.3277 -  
accuracy: 0.8680  
Epoch 81/100  
667/667 [=====] - 2s 3ms/step - loss: 0.3277 -  
accuracy: 0.8683  
Epoch 82/100  
667/667 [=====] - 2s 3ms/step - loss: 0.3274 -  
accuracy: 0.8664  
Epoch 83/100  
667/667 [=====] - 2s 3ms/step - loss: 0.3280 -  
accuracy: 0.8666  
Epoch 84/100  
667/667 [=====] - 2s 3ms/step - loss: 0.3269 -  
accuracy: 0.8673  
Epoch 85/100  
667/667 [=====] - 2s 3ms/step - loss: 0.3272 -  
accuracy: 0.8651

```

Epoch 86/100
667/667 [=====] - 2s 3ms/step - loss: 0.3277 -
accuracy: 0.8639
Epoch 87/100
667/667 [=====] - 2s 3ms/step - loss: 0.3269 -
accuracy: 0.8670
Epoch 88/100
667/667 [=====] - 2s 3ms/step - loss: 0.3268 -
accuracy: 0.8669
Epoch 89/100
667/667 [=====] - 2s 3ms/step - loss: 0.3277 -
accuracy: 0.8660
Epoch 90/100
667/667 [=====] - 2s 2ms/step - loss: 0.3272 -
accuracy: 0.8665
Epoch 91/100
667/667 [=====] - 2s 3ms/step - loss: 0.3269 -
accuracy: 0.8661
Epoch 92/100
667/667 [=====] - 2s 3ms/step - loss: 0.3278 -
accuracy: 0.8658
Epoch 93/100
667/667 [=====] - 2s 3ms/step - loss: 0.3280 -
accuracy: 0.8654
Epoch 94/100
667/667 [=====] - 2s 3ms/step - loss: 0.3274 -
accuracy: 0.8650
Epoch 95/100
667/667 [=====] - 2s 2ms/step - loss: 0.3278 -
accuracy: 0.8648
Epoch 96/100
667/667 [=====] - 2s 3ms/step - loss: 0.3270 -
accuracy: 0.8633
Epoch 97/100
667/667 [=====] - 2s 3ms/step - loss: 0.3275 -
accuracy: 0.8648
Epoch 98/100
667/667 [=====] - 2s 3ms/step - loss: 0.3272 -
accuracy: 0.8659
Epoch 99/100
667/667 [=====] - 3s 4ms/step - loss: 0.3272 -
accuracy: 0.8665
Epoch 100/100
667/667 [=====] - 3s 4ms/step - loss: 0.3274 -
accuracy: 0.8659

```

[34]: <tensorflow.python.keras.callbacks.History at 0x11dad1c2bb0>

Epoch 100/100 667/667 [=====] - 3s 4ms/step - loss: 0.3274 - accuracy: 0.8659

## 1.4 Making the predictions and evaluating the model

### 1.4.1 Predicting the result of a single observation

Therefore, our ANN model predicts that this customer stays in the bank!

**Important note 1:** Notice that the values of the features were all input in a double pair of square brackets. That’s because the “predict” method always expects a 2D array as the format of its inputs. And putting our values into a double pair of square brackets makes the input exactly a 2D array.

**Important note 2:** Notice also that the “France” country was not input as a string in the last column but as “1, 0, 0” in the first three columns. That’s because of course the predict method expects the one-hot-encoded values of the state, and as we see in the first row of the matrix of features X, “France” was encoded as “1, 0, 0”. And be careful to include these values in the first three columns, because the dummy variables are always created in the first columns.

Use our ANN model to predict if the customer with the following informations will leave the bank:

- Geography: France
- Credit Score: 600
- Gender: Male
- Age: 40 years old
- Tenure: 3 years
- Balance: \$ 60000,
- Number of Products: 2,
- Does this customer have a credit card ? Yes
- Is this customer an Active Member: Yes
- Estimated Salary: \$ 50000

So, should we say goodbye to that customer ?

- Assumption: Threshold Probability = 0.5
- Less than 0.5 = False. The person will not leave the bank
- More than 0.5 = True. The person will leave the bank

```
[39]: print(ann.predict(sc.transform([[1, 0, 0, 600, 1, 40, 3, 60000, 2, 1, 1, 50000]])) > 0.5)
```

```
[[False]]
```

```
[40]: print(ann.predict(sc.transform([[1, 0, 0, 600, 1, 40, 3, 60000, 2, 1, 1, 1, 50000]])))
```

[[0.07359219]]

- Geography: Spain  $\rightarrow [0,1,0]$
- Credit Score: 750
- Gender: Female
- Age: 35 years old

- Tenure: 10 years
- Balance: \$ 70000,
- Number of Products: 4,
- Does this customer have a credit card ? Yes
- Is this customer an Active Member: No
- Estimated Salary: \$ 50000

```
[36]: print(ann.predict(sc.transform([[0, 1, 0, 750, 1, 35, 10, 70000, 4, 1, 0,
↪50000]])) > 0.5)
```

```
[[ True]]
```

```
[41]: print(ann.predict(sc.transform([[0, 1, 0, 750, 1, 35, 10, 70000, 4, 1, 0,
↪50000]])))
```

```
[[0.8739985]]
```

### 1.4.2 Predicting the Test set results

```
[19]: y_pred = ann.predict(x_test)
y_pred = (y_pred > 0.5)
print(np.concatenate((y_pred.reshape(len(y_pred),1), y_test.
↪reshape(len(y_test),1)),1))
```

```
[[0 0]
 [0 1]
 [0 0]
 ...
 [0 0]
 [0 0]
 [0 0]]
```

### 1.4.3 Making the Confusion Matrix

```
[20]: from sklearn.metrics import confusion_matrix, accuracy_score
cm = confusion_matrix(y_test, y_pred)
print(cm)
accuracy_score(y_test, y_pred)
```


```
[[1500   95]
 [ 183  222]]
```

```
[20]: 0.861
```

```
[49]: import matplotlib.pyplot as plt
def plot_conf_mat(conf_mat):
    """
    Function that Plots a confusion matrix using seaborn's heatmap()
    """
    fig, ax = plt.subplots(figsize = (5,5))
```

```
ax = sns.heatmap(conf_mat, annot=True, cbar = False, cmap="Greens")
plt.xlabel("True label")
plt.ylabel("Predicted label");

plot_conf_mat(cm)
```



[ ]:

[ ]:

[ ]:

[ ]:

[ ]: